

Final Project Image Classification

February 6, 2024

1. Import the Required Libraries

```
[ ]: import tensorflow as tf
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.preprocessing import image
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense,
↳ Dropout
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint,
↳ ReduceLROnPlateau
from tensorflow.keras.optimizers import Adam
from tensorflow.keras.utils import plot_model
from sklearn.metrics import classification_report, confusion_matrix,
↳ ConfusionMatrixDisplay
from google.colab import files
import matplotlib.pyplot as plt
import matplotlib.image as mpimg
import numpy as np
import os
import zipfile
import time
```

2. Load Dataset

```
[ ]: !wget --no-check-certificate \
https://github.com/dicodingacademy/assets/releases/download/release/
↳ rockpaperscissors.zip \
-O /content/rockpaperscissors.zip

local_zip = '/content/rockpaperscissors.zip'
zip_ref = zipfile.ZipFile(local_zip, 'r')
zip_ref.extractall('/content')
zip_ref.close()
```

```
--2024-02-05 11:04:26-- https://github.com/dicodingacademy/assets/releases/download/release/rockpaperscissors.zip
```

```

Resolving github.com (github.com)... 20.205.243.166
Connecting to github.com (github.com)|20.205.243.166|:443... connected.
HTTP request sent, awaiting response... 302 Found
Location: https://objects.githubusercontent.com/github-production-release-
asset-2e65be/391417272/7eb836f2-695b-4a46-9c78-b65867166957?X-Amz-
Algorithm=AWS4-HMAC-SHA256&X-Amz-
Credential=AKIAVCODYLSA53PQK4ZA%2F20240205%2Fus-
east-1%2Fs3%2Faws4_request&X-Amz-Date=20240205T110426Z&X-Amz-Expires=300&X-Amz-S
ignature=846f21039a0b85251966040d86a655c0f8b0eea5f74fb8faaa4e20cfe3058921&X-Amz-
SignedHeaders=host&actor_id=0&key_id=0&repo_id=391417272&response-content-
disposition=attachment%3B%20filename%3Drockpaperscissors.zip&response-content-
type=application%2Foctet-stream [following]
--2024-02-05 11:04:26-- https://objects.githubusercontent.com/github-
production-release-
asset-2e65be/391417272/7eb836f2-695b-4a46-9c78-b65867166957?X-Amz-
Algorithm=AWS4-HMAC-SHA256&X-Amz-
Credential=AKIAVCODYLSA53PQK4ZA%2F20240205%2Fus-
east-1%2Fs3%2Faws4_request&X-Amz-Date=20240205T110426Z&X-Amz-Expires=300&X-Amz-S
ignature=846f21039a0b85251966040d86a655c0f8b0eea5f74fb8faaa4e20cfe3058921&X-Amz-
SignedHeaders=host&actor_id=0&key_id=0&repo_id=391417272&response-content-
disposition=attachment%3B%20filename%3Drockpaperscissors.zip&response-content-
type=application%2Foctet-stream
Resolving objects.githubusercontent.com (objects.githubusercontent.com)...
185.199.108.133, 185.199.109.133, 185.199.110.133, ...
Connecting to objects.githubusercontent.com
(objects.githubusercontent.com)|185.199.108.133|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 322873683 (308M) [application/octet-stream]
Saving to: '/content/rockpaperscissors.zip'

/content/rockpapers 100%[=====>] 307.92M   331MB/s   in 0.9s

2024-02-05 11:04:28 (331 MB/s) - '/content/rockpaperscissors.zip' saved
[322873683/322873683]

```

```

[ ]: base_dir = '/content/rockpaperscissors/rps-cv-images'

paper_dir = os.path.join(base_dir, 'paper')
rock_dir = os.path.join(base_dir, 'rock')
scissors_dir = os.path.join(base_dir, 'scissors')

paper_imgs = os.listdir(paper_dir)
rock_imgs = os.listdir(rock_dir)
scissors_imgs = os.listdir(scissors_dir)

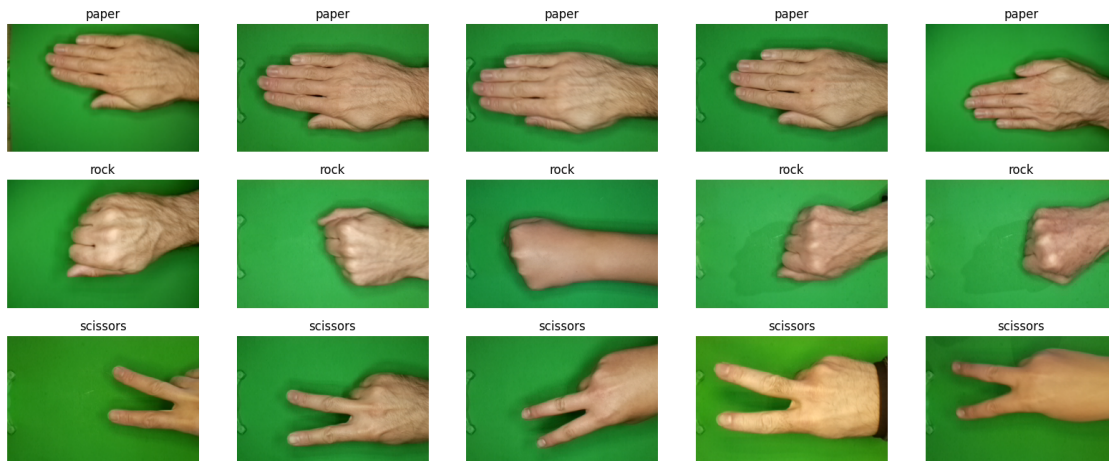
```

3. Check Image Dataset

```
[ ]: list_dir = [paper_dir, rock_dir, scissors_dir]
list_folders = [paper_imgs, rock_imgs, scissors_imgs]

plt.figure(figsize=(20, 8))
for x in range(3):
    for i, img_path in enumerate(list_folders[x][:5]):
        sp = plt.subplot(3, 5, x*5 + i + 1)
        img = mpimg.imread(os.path.join(list_dir[x], img_path))
        plt.imshow(img)
        plt.axis("off")
        plt.title(list_dir[x].split('/')[ -1])

plt.show()
```



4. Augmented Image Dataset and Split Image Dataset

```
[ ]: train_datagen = ImageDataGenerator(
    rescale=1./255,
    rotation_range=20,
    shear_range=0.2,
    horizontal_flip=True,
    fill_mode='nearest',
    validation_split=0.4)

validation_datagen = ImageDataGenerator(
    rescale=1./255,
    validation_split=0.4)
```

```
[ ]: train_generator = train_datagen.flow_from_directory(
    base_dir,
    target_size=(150, 150),
    batch_size=32,
    class_mode='categorical',
    subset='training')

validation_generator = validation_datagen.flow_from_directory(
    base_dir,
    target_size=(150, 150),
    batch_size=32,
    class_mode='categorical',
    subset='validation')
```

Found 1314 images belonging to 3 classes.
Found 874 images belonging to 3 classes.

5. Building Model

```
[ ]: model = Sequential([
    Conv2D(32, (3,3), activation='relu', padding='same', input_shape=(150, 150, 3)),
    MaxPooling2D(2,2),
    Conv2D(64, (3,3), activation='relu', padding='same'),
    MaxPooling2D(2,2),
    Conv2D(128, (3,3), activation='relu', padding='same'),
    MaxPooling2D(2,2),
    Flatten(),
    Dense(256, activation='relu'),
    Dropout(0.5),
    Dense(128, activation='relu'),
    Dropout(0.5),
    Dense(3, activation='softmax')
])
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 150, 150, 32)	896
max_pooling2d (MaxPooling2D)	(None, 75, 75, 32)	0
conv2d_1 (Conv2D)	(None, 75, 75, 64)	18496

max_pooling2d_1 (MaxPoolin g2D)	(None, 37, 37, 64)	0
conv2d_2 (Conv2D)	(None, 37, 37, 128)	73856
max_pooling2d_2 (MaxPoolin g2D)	(None, 18, 18, 128)	0
flatten (Flatten)	(None, 41472)	0
dense (Dense)	(None, 256)	10617088
dropout (Dropout)	(None, 256)	0
dense_1 (Dense)	(None, 128)	32896
dropout_1 (Dropout)	(None, 128)	0
dense_2 (Dense)	(None, 3)	387

```
=====
Total params: 10743619 (40.98 MB)
Trainable params: 10743619 (40.98 MB)
Non-trainable params: 0 (0.00 Byte)
```

```
[ ]: plot_model(
    model,
    show_shapes=True,
    show_layer_names=True,
)
```

```
[ ]:
```

conv2d_input	input:	[(None, 150, 150, 3)]
InputLayer	output:	[(None, 150, 150, 3)]



conv2d	input:	(None, 150, 150, 3)
Conv2D	output:	(None, 150, 150, 32)



max_pooling2d	input:	(None, 150, 150, 32)
MaxPooling2D	output:	(None, 75, 75, 32)



conv2d_1	input:	(None, 75, 75, 32)
Conv2D	output:	(None, 75, 75, 64)



max_pooling2d_1	input:	(None, 75, 75, 64)
MaxPooling2D	output:	(None, 37, 37, 64)



conv2d_2	input:	(None, 37, 37, 64)
Conv2D	output:	(None, 37, 37, 128)



max_pooling2d_2	input:	(None, 37, 37, 128)
MaxPooling2D	output:	(None, 18, 18, 128)



flatten	input:	(None, 18, 18, 128)
Flatten	output:	(None, 41472)



dense	input:	(None, 41472)
Dense	output:	(None, 256)



dropout	input:	(None, 256)
Dropout	output:	(None, 256)



dense_1	input:	(None, 256)
Dense	output:	(None, 128)



dropout_1	input:	(None, 128)
Dropout	output:	(None, 128)



dense_2	input:	(None, 128)
Dense	output:	(None, 3)

6. Compile Model

```
[ ]: model.compile(optimizer=Adam(),  
                  loss='categorical_crossentropy',  
                  metrics=['accuracy'])
```

7. Fit Model

```
[ ]: def train_model(model, train_generator, validation_generator, steps_per_epoch, validation_steps, epochs=40):  
    callbacks = [  
        ModelCheckpoint('best_model.h5', monitor='val_accuracy',  
        save_best_only=True, mode='max'),  
        EarlyStopping(monitor='val_loss', patience=3, mode='min'),  
        ReduceLROnPlateau(monitor='val_accuracy', patience=3, verbose=1, factor=0.  
        5, min_lr=0.000003)  
    ]  
  
    start_time = time.time()  
  
    history = model.fit(  
        train_generator,  
        steps_per_epoch=steps_per_epoch,  
        epochs=epochs,  
        validation_data=validation_generator,  
        validation_steps=validation_steps,  
        verbose=2,  
        callbacks=callbacks)  
  
    end_time = time.time()  
  
    training_time = (end_time - start_time)/60  
    print(f'Training Time: {training_time} Minute')  
  
    return history  
  
def evaluate_model(model, train_generator, validation_generator):  
    train_loss, train_acc = model.evaluate(train_generator)  
    val_loss, val_acc = model.evaluate(validation_generator)  
    if train_acc >= 0.96 and val_acc >= 0.96 and train_loss <= 0.3 and val_loss  
    <= 0.3:  
        print('The best model with the specified criteria')
```

```

        return True
    else:
        print('Not the best model with the specified criteria')
        return False

while True:
    history = train_model(
        model,
        train_generator,
        validation_generator,
        steps_per_epoch=train_generator.samples//train_generator.batch_size,
        validation_steps=validation_generator.samples//validation_generator.
↪batch_size
    )
    best_model = tf.keras.models.load_model('best_model.h5')
    if evaluate_model(best_model, train_generator, validation_generator):
        best_model.save('final_model.h5')
        break

```

Epoch 1/40

/usr/local/lib/python3.10/dist-packages/keras/src/engine/training.py:3103:

UserWarning: You are saving your model as an HDF5 file via `model.save()`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')`.

saving_api.save_model(

41/41 - 19s - loss: 1.0806 - accuracy: 0.4150 - val_loss: 0.7546 - val_accuracy: 0.7465 - lr: 0.0010 - 19s/epoch - 458ms/step

Epoch 2/40

41/41 - 12s - loss: 0.6418 - accuracy: 0.7207 - val_loss: 0.3707 - val_accuracy: 0.8877 - lr: 0.0010 - 12s/epoch - 301ms/step

Epoch 3/40

41/41 - 12s - loss: 0.4058 - accuracy: 0.8460 - val_loss: 0.2799 - val_accuracy: 0.9109 - lr: 0.0010 - 12s/epoch - 302ms/step

Epoch 4/40

41/41 - 12s - loss: 0.3446 - accuracy: 0.8690 - val_loss: 0.1865 - val_accuracy: 0.9421 - lr: 0.0010 - 12s/epoch - 300ms/step

Epoch 5/40

41/41 - 12s - loss: 0.3234 - accuracy: 0.8721 - val_loss: 0.2514 - val_accuracy: 0.9097 - lr: 0.0010 - 12s/epoch - 299ms/step

Epoch 6/40

41/41 - 13s - loss: 0.2555 - accuracy: 0.8970 - val_loss: 0.1649 - val_accuracy: 0.9583 - lr: 0.0010 - 13s/epoch - 308ms/step

Epoch 7/40

41/41 - 14s - loss: 0.2236 - accuracy: 0.9173 - val_loss: 0.1347 - val_accuracy: 0.9595 - lr: 0.0010 - 14s/epoch - 347ms/step

Epoch 8/40

41/41 - 12s - loss: 0.1901 - accuracy: 0.9337 - val_loss: 0.2264 - val_accuracy: 0.9363 - lr: 0.0010 - 12s/epoch - 296ms/step
Epoch 9/40
41/41 - 12s - loss: 0.1712 - accuracy: 0.9423 - val_loss: 0.1010 - val_accuracy: 0.9699 - lr: 0.0010 - 12s/epoch - 292ms/step
Epoch 10/40
41/41 - 12s - loss: 0.1524 - accuracy: 0.9493 - val_loss: 0.0918 - val_accuracy: 0.9711 - lr: 0.0010 - 12s/epoch - 289ms/step
Epoch 11/40
41/41 - 12s - loss: 0.1318 - accuracy: 0.9540 - val_loss: 0.0802 - val_accuracy: 0.9722 - lr: 0.0010 - 12s/epoch - 295ms/step
Epoch 12/40
41/41 - 12s - loss: 0.1240 - accuracy: 0.9602 - val_loss: 0.0792 - val_accuracy: 0.9780 - lr: 0.0010 - 12s/epoch - 298ms/step
Epoch 13/40
41/41 - 12s - loss: 0.1387 - accuracy: 0.9501 - val_loss: 0.1911 - val_accuracy: 0.9329 - lr: 0.0010 - 12s/epoch - 298ms/step
Epoch 14/40
41/41 - 12s - loss: 0.0892 - accuracy: 0.9711 - val_loss: 0.0696 - val_accuracy: 0.9780 - lr: 0.0010 - 12s/epoch - 287ms/step
Epoch 15/40

Epoch 15: ReduceLROnPlateau reducing learning rate to 0.0005000000237487257.
41/41 - 12s - loss: 0.0851 - accuracy: 0.9743 - val_loss: 0.1423 - val_accuracy: 0.9560 - lr: 0.0010 - 12s/epoch - 298ms/step
Epoch 16/40
41/41 - 12s - loss: 0.0705 - accuracy: 0.9743 - val_loss: 0.0567 - val_accuracy: 0.9873 - lr: 5.0000e-04 - 12s/epoch - 304ms/step
Epoch 17/40
41/41 - 11s - loss: 0.0551 - accuracy: 0.9883 - val_loss: 0.0565 - val_accuracy: 0.9861 - lr: 5.0000e-04 - 11s/epoch - 274ms/step
Epoch 18/40
41/41 - 12s - loss: 0.0501 - accuracy: 0.9828 - val_loss: 0.0550 - val_accuracy: 0.9861 - lr: 5.0000e-04 - 12s/epoch - 287ms/step
Epoch 19/40

Epoch 19: ReduceLROnPlateau reducing learning rate to 0.0002500000118743628.
41/41 - 12s - loss: 0.0533 - accuracy: 0.9844 - val_loss: 0.0695 - val_accuracy: 0.9873 - lr: 5.0000e-04 - 12s/epoch - 281ms/step
Epoch 20/40
41/41 - 12s - loss: 0.0415 - accuracy: 0.9875 - val_loss: 0.0649 - val_accuracy: 0.9861 - lr: 2.5000e-04 - 12s/epoch - 296ms/step
Epoch 21/40
41/41 - 13s - loss: 0.0473 - accuracy: 0.9828 - val_loss: 0.0513 - val_accuracy: 0.9896 - lr: 2.5000e-04 - 13s/epoch - 309ms/step
Epoch 22/40
41/41 - 12s - loss: 0.0404 - accuracy: 0.9836 - val_loss: 0.0589 - val_accuracy: 0.9907 - lr: 2.5000e-04 - 12s/epoch - 298ms/step

Epoch 23/40
 41/41 - 12s - loss: 0.0385 - accuracy: 0.9860 - val_loss: 0.0558 - val_accuracy: 0.9884 - lr: 2.5000e-04 - 12s/epoch - 298ms/step
 Epoch 24/40
 41/41 - 14s - loss: 0.0267 - accuracy: 0.9914 - val_loss: 0.0555 - val_accuracy: 0.9884 - lr: 2.5000e-04 - 14s/epoch - 347ms/step
 Training Time: 6.590774953365326 Minute
 42/42 [=====] - 10s 230ms/step - loss: 0.0162 - accuracy: 0.9962
 28/28 [=====] - 2s 87ms/step - loss: 0.0583 - accuracy: 0.9908
 The best model with the specified criteria

8. Plot Accuracy and Loss

```
[ ]: def plotaccloss(history):
    acc = history.history['accuracy']
    val_acc = history.history['val_accuracy']
    loss = history.history['loss']
    val_loss = history.history['val_loss']

    epochs = range(len(acc))



    plt.figure(figsize=(14, 4))

    plt.subplot(1, 2, 1)
    plt.plot(epochs, acc, 'r', label='Training accuracy')
    plt.plot(epochs, val_acc, 'b', label='Validation accuracy')
    plt.title('Training and validation accuracy')
    plt.legend(loc=0)

    plt.subplot(1, 2, 2)
    plt.plot(epochs, loss, 'r', label='Training loss')
    plt.plot(epochs, val_loss, 'b', label='Validation loss')
    plt.title('Training and validation loss')
    plt.legend(loc=0)

    plt.show()

plotaccloss(history)
```



9. Evaluate Model

```
[ ]: def evaluate(model):
    validation_generator = train_datagen.flow_from_directory(
        base_dir,
        target_size=(150, 150),
        batch_size=32,
        class_mode='categorical',
        shuffle = False,
        subset='validation')

    batch_size = 32
    num_of_test_samples = len(validation_generator.file_names)

    Y_pred = model.predict(validation_generator, num_of_test_samples //
        ↪batch_size+1)
    y_pred = np.argmax(Y_pred, axis=1)

    print('\n\nClassification Report\n')
    target_names = list(validation_generator.class_indices.keys())
    print(classification_report(validation_generator.classes, y_pred,
        ↪target_names=target_names))

    print('\nConfusion Matrix\n')
    cm = confusion_matrix(validation_generator.classes, y_pred)
    print(cm)

    disp = ConfusionMatrixDisplay(confusion_matrix=cm,
        ↪display_labels=validation_generator.class_indices.keys())
    disp.plot()
    disp.ax_.set_title("Confusion Matrix")

    evaluate(model)
```

Found 874 images belonging to 3 classes.

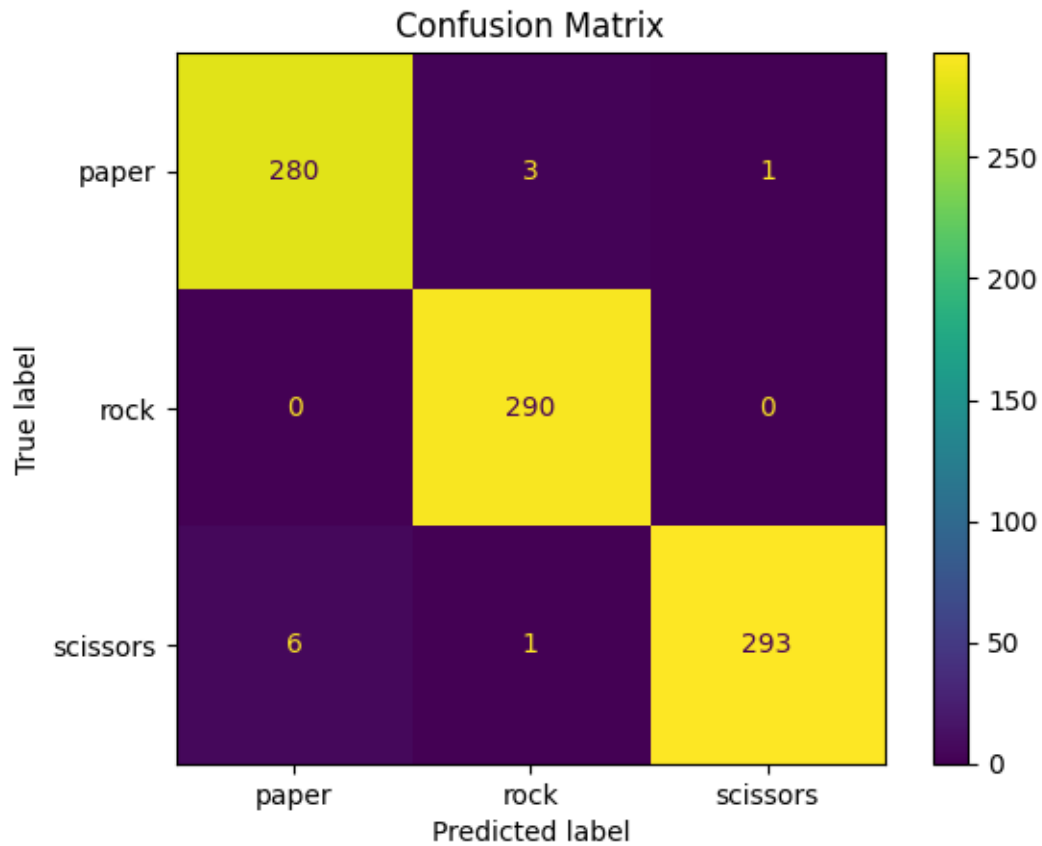
28/28 [=====] - 7s 244ms/step

Classification Report

	precision	recall	f1-score	support
paper	0.98	0.99	0.98	284
rock	0.99	1.00	0.99	290
scissors	1.00	0.98	0.99	300
accuracy			0.99	874
macro avg	0.99	0.99	0.99	874
weighted avg	0.99	0.99	0.99	874

Confusion Matrix

```
[[280  3  1]
 [ 0 290  0]
 [ 6  1 293]]
```



10. Image Prediction

```
[ ]: def imagepredict(model):
      uploaded = files.upload()

      for fn in uploaded.keys():

          path = fn
          img = image.load_img(path, target_size=(150,150))

          imgplot = plt.imshow(img)
```

```

x = image.img_to_array(img)
x = np.expand_dims(x, axis=0)
images = np.vstack([x])

classes = model.predict(images, batch_size=10)
print(fn)
if classes[0][0]>0.5:
    print(f'This is Paper!. With Probability {classes[0][0]*100}%')
elif classes[0][1]>0.5:
    print(f'This is Rock!. With Probability {classes[0][1]*100}%')
else:
    print(f'This is Scissors!. With Probability {classes[0][2]*100}%')

imagepredict(model)

```

<IPython.core.display.HTML object>

Saving scissors 1.jpeg to scissors 1.jpeg
1/1 [=====] - 0s 17ms/step
scissors 1.jpeg
This is Scissors!. With Probability 100.0%

